

Scheduling jobs on Multi Instance GPUs (MIG) using Reinforcement Learning

Tanvi Hisaria, Devyani Vij, Aayush Srivastava

ORCSE4529: Reinforcement Learning
Columbia University


December 9, 2025

Multi-Instance GPU (MIG)



MiG hardware partitioning technology is great for isolated workloads on GPUs.

Relevance of MIG Scheduling Algorithms




Maximize GPU Utilization

Cloud computing utilizes GPU scheduling

Use of AI by respondents' organizations, % of respondents

Organizations that use AI in at least 1 business function¹



Rapid AI growth demands scalable GPU capacity

Relevance of MIG Scheduling Algorithms



Energy is expensive and high in demand

Real-time AI inference workloads

- ▶ Autonomous driving
- ▶ Voice assistants
- ▶ Fraud detection
- ▶ Real-time recommendations

Real-time inference requires minimal delays.

Existing Work

Two Components of MIG Resource Management

- ▶ **Repartitioning:** Deciding how to split a GPU into MIG slices
- ▶ **Scheduling:** Assigning incoming jobs to existing slices

Prior Research

- ▶ Mao et al. (2019): RL for cluster scheduling, showing that queue-based features improve scheduling quality.
- ▶ SMART-MIG (IPDPS 2026): Mean-field MARL for *repartitioning* MIG GPUs to reduce energy and tardiness.

Scheduling Goal: Minimize Energy + Tardiness (ET)

$$ET = \frac{1}{N} \sum_{k=1}^N \frac{ae_k + t_k}{a+1},$$

where e_k is total energy, t_k is average tardiness, and a balances their importance.

MDP Formulation: SARP

State (S)

- ▶ MIG slice availability, queue features, job durations deadlines
- ▶ Added features: queue_len_norm, free_slice_fraction

Action (A)

- ▶ Assign next job to a valid MIG slice

Reward (R)

- ▶ Negative weighted combination of energy + tardiness
- ▶ Derived from ET metric: $ET = \frac{ae_k + t_k}{a+1}$

Transitions (P)

- ▶ Job executed on slice, queue updates, new jobs may arrive
- ▶ Environment evolves via simulator dynamics

Additional Setup

Environment Type

- ▶ Custom NumPy-based simulator of MIG-partitioned GPUs
- ▶ Mixed inference/training workloads with stochastic arrivals

Setting

- ▶ Simulator-based online interaction (Gym-style episodes)
- ▶ Agent learns through repeated rollouts in the simulator, not a real GPU cluster

RL Paradigm

- ▶ Policy gradient method suitable for high-dimensional, continuous state spaces
- ▶ Maskable PPO is a variant that prevents invalid actions and maintains PPO's stability, sample-efficiency, and wide use for scheduling-like tasks
PPO Clipped Objective

$$L^{\text{CLIP}}(\theta) = \mathbb{E}_t \left[\min \left(r_t(\theta) \hat{A}_t, \text{clip}(r_t(\theta), 1 - \varepsilon, 1 + \varepsilon) \hat{A}_t \right) \right]$$

$$\text{where } r_t(\theta) = \frac{\pi_\theta(a_t | s_t)}{\pi_{\theta_{\text{old}}}(a_t | s_t)}.$$



Summary of Enhancements

| Aspect | Original | Improved | Impact |
|----------------|----------------|------------------------------|--------------------------|
| Environment | Pandas | NumPy | 10–50× faster |
| Deadlines | 1.0–1.5× | 2.0–4.0× | Learnable problem |
| States Added | Features Basic | +slice sizes, +urgency, load | Better learning signal |
| Rewards | End-only | Immediate | Better credit assignment |
| Network | [256,256] | [256,256,128] | More capacity |
| Training steps | 200k | 500k | More learning |
| LR | Fixed | Annealing | Stability |
| Entropy | Fixed | Decaying | Explore → exploit |

Training Configuration (Final RL Agent)

| Parameter | Value | Improvement over Original |
|---------------|---------------|---------------------------|
| Network | [256,256,128] | Deeper (+1 layer) |
| Batch size | 4096 | 2× larger |
| Epochs | 10 | 2× more |
| Timesteps | 500,000 | 2.5× more |
| Learning rate | 3e-4 → 1e-5 | Annealing schedule |
| Entropy | 0.02 → 0.001 | Decaying entropy bonus |
| Clip range | 0.15 | Tighter updates |

Final Results: RL Vs Heuristic Baselines



Final Results: RL vs Heuristic Baselines

Performance Comparison: Enhanced RL vs Heuristics

| Method | Late Jobs (%)↓ | Avg. Tardiness↓ | Energy (MJ) |
|-------------------|----------------------------------|-----------------------------------|-----------------------------------|
| RL-PPO (Enhanced) | 37.9 ± 6.5 | 1.01 ± 0.61 | 2.44 ± 0.05 |
| EFT | 42.7 ± 5.7 | 1.02 ± 0.73 | 2.51 ± 0.05 |
| Largest-First | 42.7 ± 5.7 | 1.02 ± 0.73 | 2.51 ± 0.05 |
| Smallest-First | 54.3 ± 4.7 | 1.21 ± 0.58 | 2.42 ± 0.05 |
| Random | 50.1 ± 4.7 | 1.13 ± 0.57 | 2.49 ± 0.04 |

- ▶ RL reduces late jobs from **42.7%** (best heuristic) to **37.9%**: ~11.7% relative improvement.

Future Work

Scalability & Complexity

- ▶ Scale to larger multi-GPU clusters and more diverse workload patterns.
- ▶ Test robustness under heterogeneous GPU types and mixed job profiles.

Alternative RL Models

- ▶ Compare PPO with variants (adaptive KL, clipping strategies, GRPO).
- ▶ Explore off-policy RL (e.g., SAC) for higher sample efficiency.
- ▶ Try hierarchical RL: high-level repartitioning, low-level placement.

Joint Optimization

- ▶ Combine RL-based repartitioning with job scheduling for end-to-end optimization.
- ▶ Use bi-level control: RL selects MIG configuration; scheduler assigns jobs.

References |

- [1] Mao, H., Schwarzkopf, M., Venkatakrishnan, S. B., Meng, Z., & Alizadeh, M. (2019). *Learning scheduling algorithms for data processing clusters*. ACM SIGCOMM.
- [2] Stable-Baselines3 Documentation. "Tips and Tricks." <https://stable-baselines3.readthedocs.io/>
- [3] Andrychowicz, M., Raichuk, A., Stańczyk, P., et al. (2020). *What matters in on-policy reinforcement learning?* arXiv:2006.05990.
- [4] Schulman, J., Wolski, F., Dhariwal, P., Radford, A., & Klimov, O. (2017). *Proximal policy optimization algorithms*. arXiv:1707.06347.
- [5] RL Baselines3 Zoo. <https://github.com/DLR-RM/rl-baselines3-zoo>
- [6] Ahmed, Z., Le Roux, N., Norouzi, M., & Schuurmans, D. (2019). *Understanding the impact of entropy on policy optimization*. ICML.
- [7] <https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai/>